

Mach 1 Training Center

```
library(readxl)
Tng_Ctr_Hour2 <- read_excel("C:/RBS/Business Forecasting/Group Project/Final Project/Tng_Ctr_Hour2.xlsx")
View(Tng_Ctr_Hour2)
```

```
library(data.table)
library(ggplot2)
library(TTR)
library(fpp)
```

```
## Loading required package: forecast
```

```
## Registered S3 method overwritten by 'quantmod':
##   method      from
##   as.zoo.data.frame zoo
```

```
## Loading required package: fma
```

```
## Loading required package: expsmoother
```

```
## Loading required package: lmtest
```

```
## Loading required package: zoo
```

```
##
## Attaching package: 'zoo'
```

```
## The following objects are masked from 'package:base':
##
##   as.Date, as.Date.numeric
```

```
## Loading required package: tseries
```

```
library(fpp2)
```

```
##
## Attaching package: 'fpp2'
```

```
## The following objects are masked from 'package:fpp':
##
##   ausair, ausbeer, austa, austourists, debitcards, departures,
##   elecequip, euretail, guinearice, oil, sunspotarea, usmelec
```

```
library(ggplot2)
library(stats)
library(dplyr)
```

```
##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:data.table':
##
##   between, first, last

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library(graphics)
library(ggfortify)
```

```
## Registered S3 methods overwritten by 'ggfortify':
##   method                from
##   autoplot.Arima         forecast
##   autoplot.acf           forecast
##   autoplot.ar            forecast
##   autoplot.bats          forecast
##   autoplot.decomposed.ts forecast
##   autoplot.ets           forecast
##   autoplot.forecast      forecast
##   autoplot.stl           forecast
##   autoplot.ts            forecast
##   fitted.ar             forecast
##   fortify.ts            forecast
##   residuals.ar          forecast
```

We are evaluating the pilot training hours logged at the Mach 1 Training Center. We will create a time series forecast of the data set.

In addition to a time series forecast, we will perform multiple regression analysis using economic metrics to see if these factors may be influencing training center hours. The economic factors selected include: “Consumer Sentiment (University of Michigan)”, “NJ Unemployment Rate”, “Revenue Passenger Miles for U.S. Air Carrier Domestic and International, Scheduled Passenger Flights”, “Consumer Price Index for All Urban Consumers: All Items in U.S. City Average”, and “Median Consumer Price Index”. These economic factors were selected because they could influence the demand for flights, either directly or indirectly. This would cause increased demand for pilots and the training of pilots.

```
summary(Tng_Ctr_Hour2)
```

```
##      Year           Quarter           Month           Device_Hrs
## Length:81         Length:81         Length:81         Min.    : 222.8
## Class :character   Class :character   Class :character   1st Qu.: 899.0
## Mode  :character   Mode  :character   Mode  :character   Median :1008.0
##                                     Mean   : 990.1
##                                     3rd Qu.:1101.7
##                                     Max.   :1519.9
##
## DH_Prev_Year       DH_YoY_Change       DH_YoY_Ch_Per       Total_Inst_Hrs
## Length:81         Length:81         Length:81         Min.    : 504.6
## Class :character   Class :character   Class :character   1st Qu.:1937.3
## Mode  :character   Mode  :character   Mode  :character   Median :2203.2
##                                     Mean   :2165.7
##                                     3rd Qu.:2446.8
##                                     Max.   :3084.1
##
## Total_Inst_Hrs_Prev_Year Inst_Hrs_YoY_Change Total_Inst_Hrs_YoY_Change_Per2
## Length:81         Length:81         Length:81
## Class :character   Class :character   Class :character
## Mode  :character   Mode  :character   Mode  :character
##
##
##
## Cons_Sent          NJURN              RPM              CPIUrban
## Min.    : 70.30    Min.    : 2.900    Min.    : 2908236    Min.    :234.7
## 1st Qu.: 89.00    1st Qu.: 4.100    1st Qu.: 68459347    1st Qu.:241.2
## Median : 93.80    Median : 4.900    Median : 77115921    Median :250.8
## Mean   : 91.49    Mean   : 5.615    Mean   : 70822495    Mean   :250.2
## 3rd Qu.: 97.90    3rd Qu.: 6.200    3rd Qu.: 85326186    3rd Qu.:257.4
## Max.   :101.40    Max.   :16.600    Max.   :101794185    Max.   :274.1
##                                     NA's    :1
##
## CPIMedian
## Min.    :0.9755
## 1st Qu.:2.1551
## Median :2.5922
## Mean   :2.5862
## 3rd Qu.:2.9557
## Max.   :5.5690
##
```

Create a Factor for the Dataset

```
setDT(Tng_Ctr_Hour2)
#changing the character values into factors
Tng_Ctr_Hour2[,Quarter:=factor(Quarter)]
Tng_Ctr_Hour2[,Month:=factor(Month)]
Tng_Ctr_Hour2[,Year:=factor(Year)]
```

Create a subset for the tested data

```
FD = select(Tng_Ctr_Hour2, Year, Quarter, Month, Device_Hrs, Total_Inst_Hrs, Cons_Sent, NJURN, RPM, CPI)
```

FD

##	Year	Quarter	Month	Device_Hrs	Total_Inst_Hrs	Cons_Sent	NJURN	RPM
## 1:	2015-01	Q1	Jan	960.42	1700.67	98.1	6.8	65975447
## 2:	2015-02	Q1	Feb	944.08	1614.00	95.4	6.7	59784666
## 3:	2015-03	Q1	Mar	1429.12	2532.90	93.0	6.3	75751609
## 4:	2015-04	Q2	Apr	1097.00	2152.25	95.9	5.8	73090871
## 5:	2015-05	Q2	May	915.85	1695.43	90.7	6.0	78002935
## 6:	2015-06	Q2	Jun	783.45	1675.91	96.1	5.9	82695216
## 7:	2015-07	Q3	Jul	1034.52	2095.00	93.1	6.2	88263768
## 8:	2015-08	Q3	Aug	1169.50	2459.83	91.9	5.5	85234949
## 9:	2015-09	Q3	Sep	1027.08	2219.00	87.2	5.1	72483190
## 10:	2015-10	Q4	Oct	1262.32	2765.47	90.0	4.9	76094668
## 11:	2015-11	Q4	Nov	999.25	2239.33	91.3	4.8	70030247
## 12:	2015-12	Q4	Dec	929.42	2054.59	92.6	4.6	74827186
## 13:	2016-01	Q1	Jan	796.42	1935.51	92.0	5.3	69265865
## 14:	2016-02	Q1	Feb	874.55	2017.40	91.7	5.4	64559135
## 15:	2016-03	Q1	Mar	1091.55	2235.33	91.0	5.3	78684812
## 16:	2016-04	Q2	Apr	1141.84	2409.30	89.0	4.9	75127770
## 17:	2016-05	Q2	May	871.36	1937.34	94.7	4.8	80784690
## 18:	2016-06	Q2	Jun	1181.21	2606.56	93.5	5.0	86518457
## 19:	2016-07	Q3	Jul	757.59	1791.01	90.0	5.4	90628483
## 20:	2016-08	Q3	Aug	972.73	2216.60	89.8	5.1	86136355
## 21:	2016-09	Q3	Sep	807.02	1934.39	91.2	4.9	75323792
## 22:	2016-10	Q4	Oct	1519.92	3084.09	87.2	4.6	77274227
## 23:	2016-11	Q4	Nov	1101.67	2361.81	93.8	4.3	72215277
## 24:	2016-12	Q4	Dec	801.83	1853.99	98.2	4.3	76957615
## 25:	2017-01	Q1	Jan	995.09	2446.80	98.5	5.1	71433297
## 26:	2017-02	Q1	Feb	962.00	2169.17	96.3	5.1	64261254
## 27:	2017-03	Q1	Mar	1130.24	2768.35	96.9	4.6	80838984
## 28:	2017-04	Q2	Apr	1054.71	2291.76	97.0	4.2	79494360
## 29:	2017-05	Q2	May	1044.95	2172.54	97.1	4.3	83542041
## 30:	2017-06	Q2	Jun	1013.73	2366.74	95.0	4.5	89667877
## 31:	2017-07	Q3	Jul	693.33	1739.90	93.4	4.9	94106754
## 32:	2017-08	Q3	Aug	983.25	2304.53	96.8	4.7	90444528
## 33:	2017-09	Q3	Sep	987.64	2302.29	95.1	4.4	74645832
## 34:	2017-10	Q4	Oct	1252.69	2810.70	100.7	4.2	80629312
## 35:	2017-11	Q4	Nov	969.31	2249.47	98.5	4.2	75655182
## 36:	2017-12	Q4	Dec	806.10	1800.08	95.9	4.0	79629124
## 37:	2018-01	Q1	Jan	1060.57	2466.01	95.7	4.9	73352756
## 38:	2018-02	Q1	Feb	1200.25	2414.06	99.7	4.9	68213441
## 39:	2018-03	Q1	Mar	1262.25	2666.14	101.4	4.5	85599897
## 40:	2018-04	Q2	Apr	1184.45	2625.94	98.8	3.9	82536712
## 41:	2018-05	Q2	May	1059.92	2455.24	98.0	3.8	88019738
## 42:	2018-06	Q2	Jun	993.55	2098.89	98.2	4.2	94375329
## 43:	2018-07	Q3	Jul	908.37	1973.29	97.9	4.4	98883723
## 44:	2018-08	Q3	Aug	1096.93	2403.06	96.2	3.9	94835592
## 45:	2018-09	Q3	Sep	1121.75	2368.10	100.1	3.5	79102243

## 46:	2018-10	Q4	Oct	1412.47	2955.81	98.6	3.3	84374600
## 47:	2018-11	Q4	Nov	1010.25	2203.17	97.5	3.1	79461417
## 48:	2018-12	Q4	Dec	970.12	1991.45	98.3	3.3	82764474
## 49:	2019-01	Q1	Jan	1063.13	2542.16	91.2	4.2	76935981
## 50:	2019-02	Q1	Feb	1036.95	2441.90	93.8	4.0	70762598
## 51:	2019-03	Q1	Mar	1130.87	2456.02	98.4	3.6	90028517
## 52:	2019-04	Q2	Apr	903.97	2286.02	97.2	2.9	86209795
## 53:	2019-05	Q2	May	1284.95	2734.56	100.0	3.0	92550535
## 54:	2019-06	Q2	Jun	1265.56	2571.35	98.2	3.2	97811522
## 55:	2019-07	Q3	Jul	848.64	2075.30	98.4	3.8	101794185
## 56:	2019-08	Q3	Aug	1247.40	2767.26	89.8	3.5	98025331
## 57:	2019-09	Q3	Sep	1106.84	2441.50	93.2	3.2	83345979
## 58:	2019-10	Q4	Oct	1217.08	2626.36	95.5	3.2	87646386
## 59:	2019-11	Q4	Nov	1091.84	2377.05	96.8	3.2	80521747
## 60:	2019-12	Q4	Dec	1024.67	2085.33	99.3	3.2	89959899
## 61:	2020-01	Q1	Jan	1094.62	2523.89	99.8	4.1	81046955
## 62:	2020-02	Q1	Feb	1050.98	2137.86	101.0	3.9	73892509
## 63:	2020-03	Q1	Mar	726.19	1556.44	89.1	4.2	42770967
## 64:	2020-04	Q2	Apr	222.80	504.57	71.8	16.4	2908236
## 65:	2020-05	Q2	May	556.92	1181.00	72.3	16.6	7137356
## 66:	2020-06	Q2	Jun	899.00	1831.79	78.1	16.1	14806854
## 67:	2020-07	Q3	Jul	585.58	1427.42	72.5	13.6	22062675
## 68:	2020-08	Q3	Aug	811.74	1982.89	74.1	10.9	23569064
## 69:	2020-09	Q3	Sep	1047.41	2283.34	80.4	7.7	23073723
## 70:	2020-10	Q4	Oct	1239.26	2568.26	81.8	7.1	27870756
## 71:	2020-11	Q4	Nov	911.93	1968.93	76.9	9.6	27957384
## 72:	2020-12	Q4	Dec	569.75	1303.50	80.7	7.2	30832402
## 73:	2021-01	Q1	Jan	685.91	1685.08	79.0	8.0	27800390
## 74:	2021-02	Q1	Feb	692.88	1605.12	76.8	8.2	26323164
## 75:	2021-03	Q1	Mar	805.42	1810.00	84.9	7.8	42683735
## 76:	2021-04	Q2	Apr	904.00	2178.17	88.3	7.1	47644846
## 77:	2021-05	Q2	May	937.62	1977.58	82.9	7.0	57822304
## 78:	2021-06	Q2	Jun	954.00	2056.29	85.5	7.9	68541316
## 79:	2021-07	Q3	Jul	605.00	1457.42	81.2	7.5	78168642
## 80:	2021-08	Q3	Aug	1027.23	2175.39	70.3	6.9	71714174
## 81:	2021-09	Q3	Sep	1008.00	2173.00	72.8	6.2	NA
##	Year	Quarter	Month	Device_Hrs	Total_Inst_Hrs	Cons_Sent	NJURN	RPM
##	CPIUrban	CPIMedian						
## 1:	234.747	1.9475296						
## 2:	235.342	1.9544945						
## 3:	235.976	2.4333359						
## 4:	236.222	2.9626505						
## 5:	237.001	2.5012911						
## 6:	237.657	2.9723996						
## 7:	238.034	1.9220253						
## 8:	238.033	1.6357906						
## 9:	237.498	2.9984157						
## 10:	237.733	2.2831645						
## 11:	238.017	2.2654806						
## 12:	237.761	1.8347197						
## 13:	237.652	2.7084166						
## 14:	237.336	2.5377144						
## 15:	238.080	2.2220046						
## 16:	238.992	3.8254779						

17: 239.557 3.2805192
18: 240.222 2.1033905
19: 240.101 2.4901647
20: 240.545 2.7425709
21: 241.176 2.4983293
22: 241.741 2.0180256
23: 242.026 2.4588315
24: 242.637 2.0383851
25: 243.620 3.3492298
26: 243.872 2.6177801
27: 243.766 1.4650202
28: 244.274 2.0820924
29: 244.069 2.1551132
30: 244.218 1.8236004
31: 244.280 2.1967409
32: 245.205 2.7462102
33: 246.551 2.8264339
34: 246.657 2.3877066
35: 247.378 2.2123387
36: 247.736 2.9557472
37: 248.721 3.6292773
38: 249.300 2.0689953
39: 249.517 2.2785928
40: 250.275 2.9320091
41: 250.786 3.1113162
42: 251.152 3.0118933
43: 251.345 2.2764981
44: 251.735 2.2967113
45: 252.183 2.6782871
46: 252.899 2.3248663
47: 252.822 3.4834371
48: 252.493 2.9268951
49: 252.441 2.9421529
50: 252.969 2.8563444
51: 254.147 2.9894583
52: 255.326 3.1675885
53: 255.371 2.5585351
54: 255.423 3.9747278
55: 255.925 2.6178534
56: 256.118 2.6257566
57: 256.532 3.1506552
58: 257.387 2.7022755
59: 257.989 2.9606861
60: 258.203 1.9992224
61: 258.687 2.5922152
62: 258.824 2.7887863
63: 257.989 2.5767300
64: 256.192 1.7631315
65: 255.942 3.0919936
66: 257.282 1.7091233
67: 258.604 2.8424264
68: 259.511 2.4333628
69: 260.149 1.1658029
70: 260.462 2.7957567

```
## 71: 260.927 0.9755080
## 72: 261.560 1.7310060
## 73: 262.231 0.9761601
## 74: 263.161 2.8363413
## 75: 264.793 1.8331691
## 76: 266.832 2.8953414
## 77: 268.551 3.2082219
## 78: 270.981 2.9056238
## 79: 272.265 3.6798008
## 80: 273.012 4.0950754
## 81: 274.138 5.5690464
##      CPIUrban CPIMedian
```

```
summary(FD)
```

```
##      Year      Quarter      Month      Device_Hrs      Total_Inst_Hrs
## 2015-01: 1    Q1:21    Apr      : 7    Min.      : 222.8    Min.      : 504.6
## 2015-02: 1    Q2:21    Aug      : 7    1st Qu.: 899.0    1st Qu.:1937.3
## 2015-03: 1    Q3:21    Feb      : 7    Median :1008.0    Median :2203.2
## 2015-04: 1    Q4:18    Jan      : 7    Mean   : 990.1    Mean   :2165.7
## 2015-05: 1                Jul      : 7    3rd Qu.:1101.7    3rd Qu.:2446.8
## 2015-06: 1                Jun      : 7    Max.   :1519.9    Max.   :3084.1
## (Other):75                (Other):39
##      Cons_Sent      NJURN      RPM      CPIUrban
## Min.      : 70.30    Min.      : 2.900    Min.      : 2908236    Min.      :234.7
## 1st Qu.: 89.00    1st Qu.: 4.100    1st Qu.: 68459347    1st Qu.:241.2
## Median : 93.80    Median : 4.900    Median : 77115921    Median :250.8
## Mean   : 91.49    Mean   : 5.615    Mean   : 70822495    Mean   :250.2
## 3rd Qu.: 97.90    3rd Qu.: 6.200    3rd Qu.: 85326186    3rd Qu.:257.4
## Max.   :101.40    Max.   :16.600    Max.   :101794185    Max.   :274.1
##                                     NA's      :1
##      CPIMedian
## Min.      :0.9755
## 1st Qu.:2.1551
## Median :2.5922
## Mean   :2.5862
## 3rd Qu.:2.9557
## Max.   :5.5690
##
```

Convert to Time Series Data


```
FDts = ts(FD$Device_Hrs, frequency = 12, start = c(2015,1))
FDts
```

```
##      Jan      Feb      Mar      Apr      May      Jun      Jul      Aug      Sep
## 2015  960.42  944.08 1429.12 1097.00  915.85  783.45 1034.52 1169.50 1027.08
## 2016  796.42  874.55 1091.55 1141.84  871.36 1181.21  757.59  972.73  807.02
## 2017  995.09  962.00 1130.24 1054.71 1044.95 1013.73  693.33  983.25  987.64
## 2018 1060.57 1200.25 1262.25 1184.45 1059.92  993.55  908.37 1096.93 1121.75
```

```
## 2019 1063.13 1036.95 1130.87 903.97 1284.95 1265.56 848.64 1247.40 1106.84
## 2020 1094.62 1050.98 726.19 222.80 556.92 899.00 585.58 811.74 1047.41
## 2021 685.91 692.88 805.42 904.00 937.62 954.00 605.00 1027.23 1008.00
##      Oct      Nov      Dec
## 2015 1262.32 999.25 929.42
## 2016 1519.92 1101.67 801.83
## 2017 1252.69 969.31 806.10
## 2018 1412.47 1010.25 970.12
## 2019 1217.08 1091.84 1024.67
## 2020 1239.26 911.93 569.75
## 2021
```

Create a plot of the time series

```
plot(FDts)
```



The plot shows a series of peaks and valleys, which suggests seasonality. Hours seemed to flow between ~700 and ~1500 consistently until 2020 when COVID struck. This caused a drop to 222 hrs in April. October is the most busy month, never falling below 1200 hrs. July and December appear to be the lightest months, only surpassing 1000 hours once each during the observed history.

Because of the pandemic, we will remove the data from 2020 onward as it was influenced by the COVID pandemic, rendering the observations afterwards as inconsistent with the forecast


Training the Model

```
FD_train = ts(FD$Device_Hrs, frequency = 12, start = c(2015,1), end = c(2018, 12))
```

```
FD_test = ts(FD$Device_Hrs, frequency = 12, start = c(2019,1), end = c(2019, 12))
```

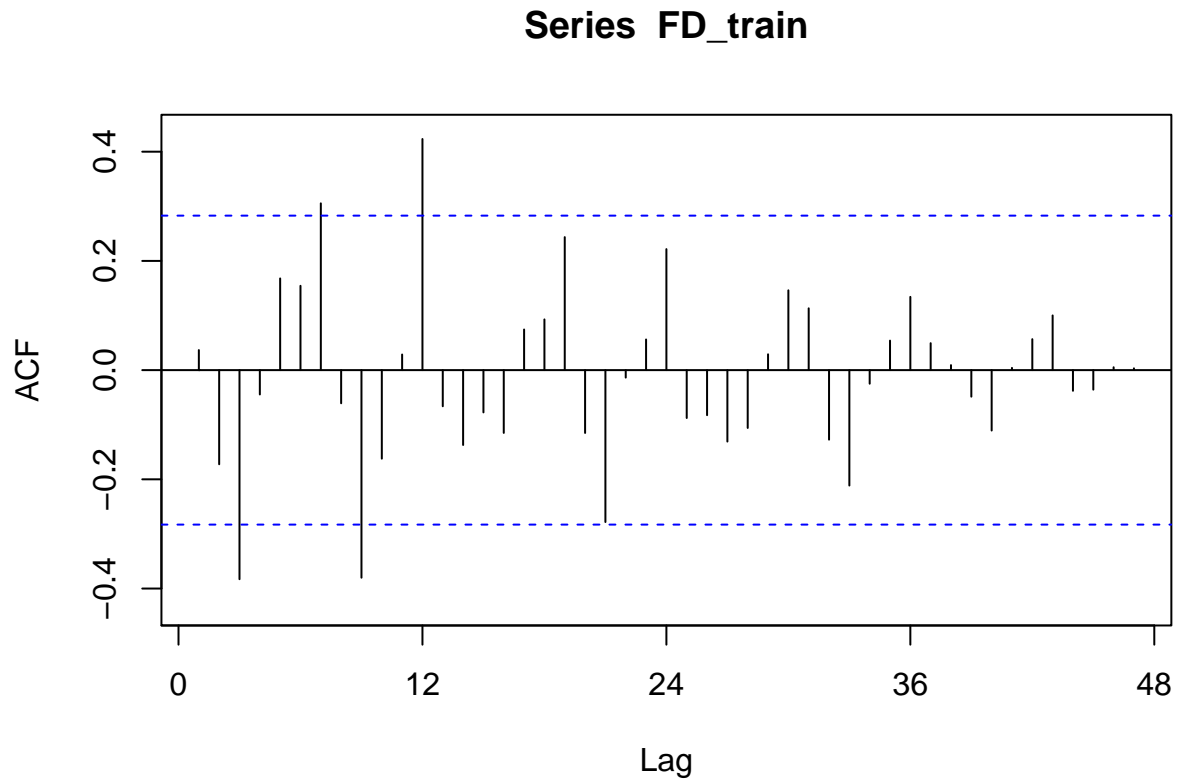
Autocorrelation of the data

```
Acf(FD_train)
```



We observe strong positive and negative autocorrelation, which furthers our suspicions that there is seasonality

```
Acf(FD_train, lag.max = 48)
```



```
# Checking For Seasonality and Trends
```

```
fit= stl(FD_train, s.window = 5)  
fit
```


```
## Call:  
## stl(x = FD_train, s.window = 5)  
##  
## Components  
##      seasonal      trend  remainder  
## Jan 2015 -129.045449 1072.110   17.355668  
## Feb 2015 -107.262914 1066.307  -14.963741  
## Mar 2015  224.832060 1060.504  143.784413  
## Apr 2015   88.044656 1054.700  -45.745057  
## May 2015 -100.010890 1049.299  -33.438306  
## Jun 2015  -38.452890 1043.898 -221.995102  
## Jul 2015 -142.974820 1038.497  138.998032  
## Aug 2015   49.132423 1033.557   86.811014  
## Sep 2015  -78.637347 1028.616   77.101009  
## Oct 2015  352.636565 1023.676 -113.992678  
## Nov 2015   25.472369 1019.294  -45.516306  
## Dec 2015 -149.739347 1014.912   64.247586  
## Jan 2016 -117.920199 1010.530  -96.189386  
## Feb 2016  -96.644317 1009.011  -37.816415
```

```

## Mar 2016 185.174945 1007.492 -101.116824
## Apr 2016 91.762657 1005.973 44.104317
## May 2016 -78.120110 1005.592 -56.111737
## Jun 2016 7.703243 1005.211 168.296088
## Jul 2016 -192.946834 1004.829 -54.292656
## Aug 2016 19.761423 1005.800 -52.831165
## Sep 2016 -91.470378 1006.770 -108.279615
## Oct 2016 355.701025 1007.740 156.478730
## Nov 2016 19.348695 1008.064 74.257790
## Dec 2016 -176.476142 1008.387 -30.080643
## Jan 2017 -60.060469 1008.710 46.440415
## Feb 2017 -13.295386 1008.911 -33.615230
## Mar 2017 137.033245 1009.111 -15.904423
## Apr 2017 95.565533 1009.312 -50.167271
## May 2017 -24.715705 1012.262 57.403293
## Jun 2017 28.641174 1015.213 -30.124259
## Jul 2017 -256.977073 1018.164 -67.856686
## Aug 2017 -21.654380 1022.593 -17.688292
## Sep 2017 -64.283047 1027.022 24.901463
## Oct 2017 332.006351 1031.450 -110.766847
## Nov 2017 -28.975050 1038.667 -40.382206
## Dec 2017 -201.159951 1045.884 -38.624064
## Jan 2018 -29.025710 1053.101 36.494935
## Feb 2018 29.202447 1061.037 110.010983
## Mar 2018 154.571226 1068.972 38.706409
## Apr 2018 94.558383 1076.908 12.983457
## May 2018 -11.493527 1085.086 -13.672874
## Jun 2018 -9.889678 1093.265 -89.824965
## Jul 2018 -245.915270 1101.443 52.842385
## Aug 2018 -18.182053 1109.607 5.505051
## Sep 2018 -36.462252 1117.771 40.441132
## Oct 2018 304.327121 1125.935 -17.792358
## Nov 2018 -59.028409 1133.988 -64.710038
## Dec 2018 -194.556831 1142.042 22.635174

```

```
plot(fit)
```




We can see some seasonality though it is not perfect. The trend is interesting as it appears that training hours were trending down, only to rebound.

Forecasting Data

Simple forecast of three periods

```
FitFore3 = forecast(fit, h=3)
plot(FitFore3)
```

Forecasts from STL + ETS(A,N,N)




Over the next three periods, we expect the number of hours to increase

What about 6 periods?

```
FitFore6 = forecast(fit, h=6)
plot(FitFore6)
```

Forecasts from STL + ETS(A,N,N)




We expect the data to increase, then decrease after three periods, consistent with our history

What about 12 Periods?!?!

```
FitFore12 = forecast(fit, h=12)
plot(FitFore12)
```

Forecasts from STL + ETS(A,N,N)



This is interesting and shows the seasonality of the data. It doesn't appear to be an exact duplication of the previous 12 months.

```
accuracy(FitFore3)
```

Accuracy Test of Forecast

```
##               ME      RMSE      MAE      MPE      MAPE      MASE
## Training set  7.149583 88.13674 72.24431 -0.02977668 7.049975 0.5119645
##               ACF1
## Training set -0.04022578
```

```
accuracy(FitFore6)
```

```
##               ME      RMSE      MAE      MPE      MAPE      MASE
## Training set  7.149583 88.13674 72.24431 -0.02977668 7.049975 0.5119645
##               ACF1
## Training set -0.04022578
```

```
accuracy(FitFore12)
```

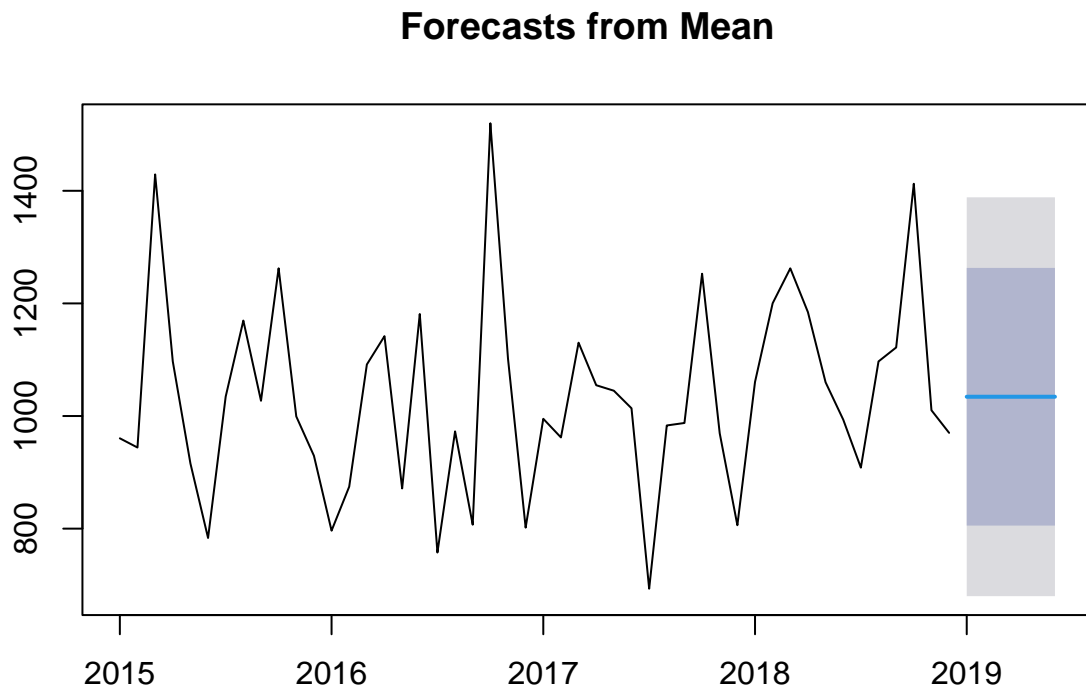
```
##               ME      RMSE      MAE      MPE      MAPE      MASE
```

```
## Training set 7.149583 88.13674 72.24431 -0.02977668 7.049975 0.5119645
##               ACF1
## Training set -0.04022578
```

All forecasts have the same accuracy measures. Length of the forecast does not impact the accuracy. Interesting.

Mean Forecast Method


```
mean_FDT <- meanf(FD_train,6) # 6 is the forecasting period (6 quarters out)
plot(mean_FDT)
```



12 Month Forecast of Trained data (raw) vs. FIT data


```
Forecast_Train12 = forecast(FD_train, h=12)
plot(Forecast_Train12)
```


Forecasts from ETS(M,N,M)



```
FitFore12 = forecast(fit, h=12)  
plot(FitFore12)
```

Forecasts from STL + ETS(A,N,N)



Fit forecast is steeper. Is it capturing more of the seasonal swings?

```
accuracy(FitFore12)
```

Let's compare the accuracy of the forecasts

```
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set  7.149583 88.13674 72.24431 -0.02977668 7.049975 0.5119645
##              ACF1
## Training set -0.04022578
```

```
accuracy(Forecast_Train12)
```


```
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set  8.442611 105.9033 87.19015 -0.2842957 8.633811 0.6178793
##              ACF1
## Training set  0.02967374
```

Fit forecast appears more accurate. Has lower RMSE and MAPE.

Naive Forecast

```
naive_forecast <- naive(FD_train,12)  
plot(naive_forecast)
```


Forecasts from Naive method



Not particularly useful because the the scope stretches into negative territory. Lets trim down the forecast to something shorter

```
naive_forecast6 <- naive(FD_train,6)  
plot(naive_forecast6)
```


Forecasts from Naive method



Still not good. 3 Months?

```
naive_forecast3 <- naive(FD_train,3)
plot(naive_forecast3)
```

Forecasts from Naive method




The range of outcomes still appears too large. Naive doesn't look like the best model.

Random Walk Forecast

```
rwf_forecast <- rwf(FD_train,6)
plot(rwf_forecast)
```

Forecasts from Random walk



This graph also shows little by way of seasonality.

Lets compare the accuracy of all the forecasts:

```
accuracy(Forecast_Train12)
```

```
##           ME      RMSE      MAE      MPE      MAPE      MASE
## Training set  8.442611 105.9033  87.19015 -0.2842957  8.633811  0.6178793
##           ACF1
## Training set  0.02967374
```

```
accuracy(FitFore12)
```

```
##           ME      RMSE      MAE      MPE      MAPE      MASE
## Training set  7.149583  88.13674  72.24431 -0.02977668  7.049975  0.5119645
##           ACF1
## Training set -0.04022578
```

```
accuracy(rwf_forecast)
```

```
##           ME      RMSE      MAE      MPE      MAPE      MASE      ACF1
## Training set  0.206383 241.3982 194.2936 -2.591405 18.87992 1.376876 -0.3923039
```




Fit forecast appears the strongest based on these metrics. It also captures the seasonality well.

Check Residuals for Fit Forecast

```
checkresiduals(FitFore12)
```

```
## Warning in checkresiduals(FitFore12): The fitted degrees of freedom is based on  
## the model used for the seasonally adjusted data.
```

Residuals from STL + ETS(A,N,N)



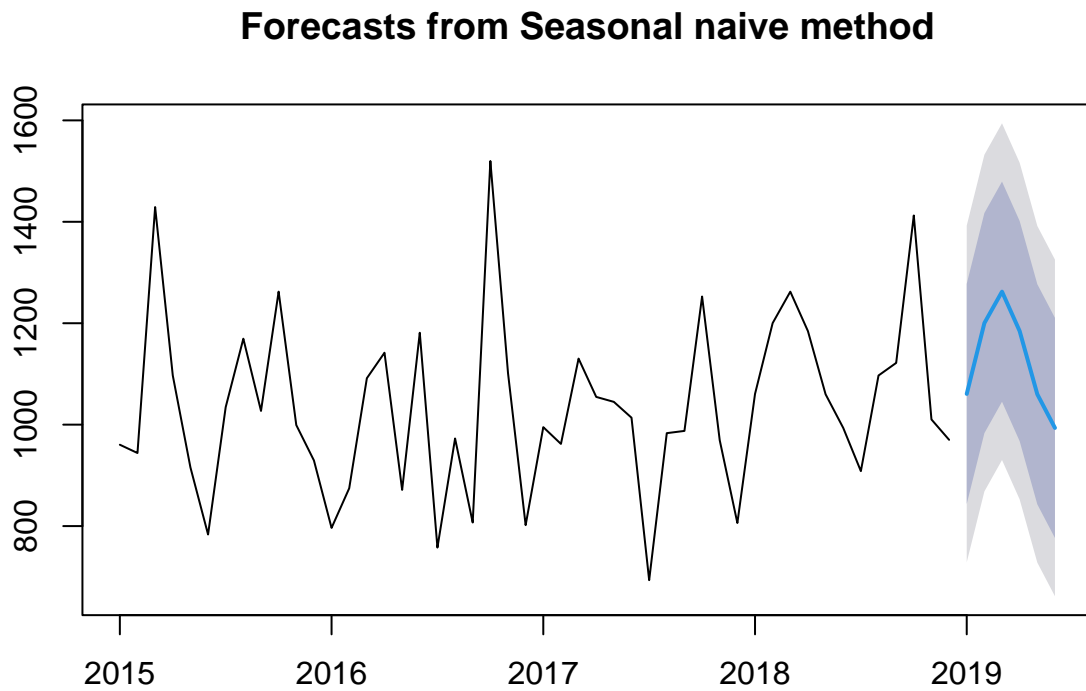
```
##  
## Ljung-Box test  
##  
## data: Residuals from STL + ETS(A,N,N)  
## Q* = 10.05, df = 8, p-value = 0.2615  
##  
## Model df: 2. Total lags used: 10
```

Residuals appear to be normally distributed with some slight skewing to the left. P-values are above .05 meaning

what? ^^^


Seasonal Naive

```
snaive_forecast <- snaive(FD_train, 6)
plot(snaive_forecast)
```



```
snaive_Fit <- snaive(FitFore6, 6)
plot(snaive_Fit)
```


Forecasts from Seasonal naive method



These results are interesting and very different.

```
# moving averages
```

```
MA5_forecast <- ma(FD_train ,order=6)
```

```
plot(FD_train) +
```

```
lines(MA5_forecast,col="Green")
```


```
## integer(0)
```

```
MA9_forecast <- ma(FD_train,order=9)
```

```
lines(MA9_forecast,col="Blue")
```

```
MA50_forecast <- ma(FD_train,order=10, centre = FALSE)
```


```
lines(MA50_forecast,col="Purple")
```



Holt Winters

```
HW_FDTrain = HoltWinters(FD_train)
plot(HW_FDTrain)
```

Holt-Winters filtering



```
HW_FDTrain
```


```
## Holt-Winters exponential smoothing with trend and additive seasonal component.
##
## Call:
## HoltWinters(x = FD_train)
##
## Smoothing parameters:
##  alpha: 0.01670918
##  beta : 1
##  gamma: 0.7138949
##
## Coefficients:
##           [,1]
## a    1090.033935
## b      12.085411
## s1     38.103050
## s2    134.598709
## s3    226.239016
## s4    153.356954
## s5     32.759172
## s6     -1.165754
## s7   -163.195873
## s8     39.529001
## s9     36.894803
## s10    337.423704
```

```
## s11 -45.657127
## s12 -135.072455
```

```
HW_FDTrain2 = HoltWinters(x = FD_train, beta = FALSE, gamma = FALSE)
HW_FDTrain2
```


```
## Holt-Winters exponential smoothing without trend and without seasonal component.
##
## Call:
## HoltWinters(x = FD_train, beta = FALSE, gamma = FALSE)
##
## Smoothing parameters:
## alpha: 0.04889481
## beta : FALSE
## gamma: FALSE
##
## Coefficients:
##      [,1]
## a 1042.894
```

```
plot(HW_FDTrain2)
```



```
SSE_Simple <- HoltWinters(FD_train, beta=FALSE, gamma=FALSE)
plot(SSE_Simple)
```


Holt-Winters filtering



SSE with Trend but no seasonality so $\gamma = \text{FALSE}$

```
SSE_Trend <- HoltWinters(FD_train,gamma=FALSE)
plot(SSE_Trend)
```

Holt-Winters filtering



```
SSE_Simple$SSE
```

```
## [1] 1534136
```

```
SSE_Trend$SSE
```

```
## [1] 1617916
```

ETS

```
ets_forecast <- ets(FD_train)
attributes(ets_forecast)
```

```
## $names
## [1] "loglik"      "aic"         "bic"         "aicc"        "mse"
## [6] "amse"        "fit"         "residuals"   "fitted"      "states"
## [11] "par"         "m"           "method"      "series"      "components"
## [16] "call"        "initstate"   "sigma2"      "x"
##
## $class
## [1] "ets"
```


```
ets_forecast$mse
```

```
## [1] 11215.5
```


```
ets_forecast
```

```
## ETS(M,N,M)
##
## Call:
## ets(y = FD_train)
##
## Smoothing parameters:
##   alpha = 0.0523
##   gamma = 0.0012
##
## Initial states:
##   l = 1032.3831
##   s = 0.8324 0.9804 1.3176 0.9464 1.0033 0.8335
##       0.9825 0.9549 1.078 1.1775 0.9744 0.9192
##
## sigma: 0.1249
##
##      AIC      AICc      BIC
## 664.3594 679.3594 692.4274
```

```
autoplot(ets_forecast)
```




```
cbind('Residuals' = residuals(ets_forecast),
      'Forecast errors' = residuals(ets_forecast,type='response')) %>%
autoplot(facet=TRUE) + xlab("Year") + ylab("")
```

```
checkresiduals(ets_forecast)
```

Residuals from ETS(M,N,M)



```
##
##  Ljung-Box test
##
## data:  Residuals from ETS(M,N,M)
## Q* = 20.465, df = 3, p-value = 0.000136
##
## Model df: 14.   Total lags used: 17
```

Residuals are skewed right, not the most normal we've seen so far.

Forecast with ETS

```
# forecast with ets
forecast.ets(ets_forecast, h=6)
```


	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Jan 2019	968.184	813.2283	1123.140	731.1997	1205.168
## Feb 2019	1026.268	861.7881	1190.748	774.7176	1277.818
## Mar 2019	1240.202	1041.1600	1439.245	935.7933	1544.611
## Apr 2019	1135.404	952.9295	1317.879	856.3332	1414.475
## May 2019	1005.668	843.8214	1167.515	758.1447	1253.192
## Jun 2019	1034.692	867.9455	1201.439	779.6752	1289.709

```

forecast_ets <- forecast.ets(ets_forecast, h=6)
plot(forecast_ets)
lines(forecast_ets$fitted, col = "lightblue")
lines(FD_test, col = 'orange')

```

Forecasts from ETS(M,N,M)



The forecast shows the forecasted points fall within the confidence intervals

```
accuracy(forecast_ets)
```

```

##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 8.442611 105.9033 87.19015 -0.2842957 8.633811 0.6178793
##              ACF1
## Training set 0.02967374

```

```
accuracy(Forecast_Train12)
```

```

##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 8.442611 105.9033 87.19015 -0.2842957 8.633811 0.6178793
##              ACF1
## Training set 0.02967374

```

```
accuracy(FitFore12)
```

```

##              ME      RMSE      MAE      MPE      MAPE      MASE

```

```
## Training set 7.149583 88.13674 72.24431 -0.02977668 7.049975 0.5119645
##               ACF1
## Training set -0.04022578
```

```
accuracy(rwf_forecast)
```

```
##               ME      RMSE      MAE      MPE      MAPE      MASE      ACF1
## Training set 0.206383 241.3982 194.2936 -2.591405 18.87992 1.376876 -0.3923039
```

ETS Forecast is the same as the fit forecast?


Decomposition

```
stl_decomp <- stl(FD_train,s.window = 'periodic')
stl_decomp
```

```
## Call:
## stl(x = FD_train, s.window = "periodic")
##
## Components
##      seasonal      trend      remainder
## Jan 2015  -80.24561 1062.4550 -21.789419
## Feb 2015  -37.81585 1058.0382 -76.142381
## Mar 2015  195.58890 1053.6214 179.909677
## Apr 2015   88.16603 1048.6299 -39.795893
## May 2015  -56.94693 1043.6383 -70.841376
## Jun 2015  -37.66431 1037.6384 -216.524096
## Jul 2015 -182.87953 1031.6385 185.761019
## Aug 2015   21.95281 1025.0648 122.482424
## Sep 2015  -50.09491 1018.4910  58.683883
## Oct 2015  324.02207 1013.0154 -74.717425
## Nov 2015  -19.56858 1007.5397  11.278884
## Dec 2015 -164.51421 1004.8528  89.081460
## Jan 2016  -80.24561 1002.1658 -125.500201
## Feb 2016  -37.81585  996.0810 -83.715141
## Mar 2016  195.58890  989.9962 -94.035062
## Apr 2016   88.16603  991.1921  62.481865
## May 2016  -56.94693  992.3881 -64.081122
## Jun 2016  -37.66431  999.3013 219.573013
## Jul 2016 -182.87953 1006.2145 -65.745018
## Aug 2016   21.95281 1011.9359 -61.158694
## Sep 2016  -50.09491 1017.6572 -160.542316
## Oct 2016  324.02207 1018.6284 177.269570
## Nov 2016  -19.56858 1019.5995 101.639074
## Dec 2016 -164.51421 1018.0979 -51.753707
## Jan 2017  -80.24561 1016.5963  58.739275
## Feb 2017  -37.81585 1014.0360 -14.220179
## Mar 2017  195.58890 1011.4757 -76.824613
## Apr 2017   88.16603 1005.3158 -38.771854
## May 2017  -56.94693  999.1559 102.740991
## Jun 2017  -37.66431  997.9515  53.442775
```

```
## Jul 2017 -182.87953 996.7471 -120.537606
## Aug 2017 21.95281 1004.9838 -43.686579
## Sep 2017 -50.09491 1013.2204 24.514502
## Oct 2017 324.02207 1023.7274 -95.059453
## Nov 2017 -19.56858 1034.2344 -45.355789
## Dec 2017 -164.51421 1043.6364 -73.022142
## Jan 2018 -80.24561 1053.0383 87.777269
## Feb 2018 -37.81585 1064.5129 173.552900
## Mar 2018 195.58890 1075.9876 -9.326450
## Apr 2018 88.16603 1080.8642 15.419734
## May 2018 -56.94693 1085.7409 31.126005
## Jun 2018 -37.66431 1089.6332 -58.418844
## Jul 2018 -182.87953 1093.5254 -2.275859
## Aug 2018 21.95281 1096.4897 -21.512523
## Sep 2018 -50.09491 1099.4540 72.390867
## Oct 2018 324.02207 1101.7365 -13.288588
## Nov 2018 -19.56858 1104.0190 -74.200425
## Dec 2018 -164.51421 1105.8608 28.773428
```

```
plot(stl_decomp)
```



Again the data shows seasonality


Table of seasonal adjustment

```
SA_ts <- seasadj(stl_decomp)
SA_ts
```

##		Jan	Feb	Mar	Apr	May	Jun	Jul
##	2015	1040.6656	981.8958	1233.5311	1008.8340	972.7969	821.1143	1217.3995
##	2016	876.6656	912.3658	895.9611	1053.6740	928.3069	1218.8743	940.4695
##	2017	1075.3356	999.8158	934.6511	966.5440	1101.8969	1051.3943	876.2095
##	2018	1140.8156	1238.0658	1066.6611	1096.2840	1116.8669	1031.2143	1091.2495
##		Aug	Sep	Oct	Nov	Dec		
##	2015	1147.5472	1077.1749	938.2979	1018.8186	1093.9342		
##	2016	950.7772	857.1149	1195.8979	1121.2386	966.3442		
##	2017	961.2972	1037.7349	928.6679	988.8786	970.6142		
##	2018	1074.9772	1171.8449	1088.4479	1029.8186	1134.6342		

It's hard to draw conclusions with the data in this format.

```
plot(SA_ts)
```



The Data looks even more volatile when you adjust for seasonality


Forecasting the decomposition is typically not very useful, but we will try it.

```
# forecast
f_stl <- forecast(stl_decomp, h=5)
f_stl
```

```
##          Point Forecast      Lo 80      Hi 80      Lo 95      Hi 95
## Jan 2019         954.0297  818.4401 1089.619  746.6632 1161.396
## Feb 2019         996.4595  860.8698 1132.049  789.0930 1203.826
## Mar 2019        1229.8642 1094.2746 1365.454 1022.4977 1437.231
## Apr 2019        1122.4414  986.8517 1258.031  915.0749 1329.808
## May 2019         977.3284  841.7387 1112.918  769.9619 1184.695
```

```
plot(f_stl)
lines(f_stl$fitted, col = "lightblue")
lines(FD_test, col = 'red')
```

Forecasts from STL + ETS(M,N,N)



```
accuracy(f_stl)
```

```
##          ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -0.00529333 103.5708 86.01888 -1.071521 8.579715 0.609579
##          ACF1
## Training set 0.05904374
```

```
accuracy(ets_forecast)
```

```
##           ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 8.442611 105.9033 87.19015 -0.2842957 8.633811 0.6178793
##           ACF1
## Training set 0.02967374
```

This appears to be the strongest forecast model we have found. Lower MAPE and RMSE

Multiple Regression Analysis

```
setDT(Tng_Ctr_Hour2)
class(Tng_Ctr_Hour2)
```

```
## [1] "data.table" "data.frame"
```

```
Tng_Ctr_Hour2[, Quarter := factor(Quarter, ordered = T)]
Tng_Ctr_Hour2[, Month := factor(Month, ordered = T)]
summary(Tng_Ctr_Hour2)
```

```
##      Year  Quarter  Month  Device_Hrs  DH_Prev_Year
## 2015-01: 1  Q1:21  Apr    : 7  Min.    : 222.8  Length:81
## 2015-02: 1  Q2:21  Aug    : 7  1st Qu.: 899.0  Class :character
## 2015-03: 1  Q3:21  Feb    : 7  Median :1008.0  Mode  :character
## 2015-04: 1  Q4:18  Jan    : 7  Mean    : 990.1
## 2015-05: 1      Jul  : 7  3rd Qu.:1101.7
## 2015-06: 1      Jun  : 7  Max.    :1519.9
## (Other):75      (Other):39
## DH_YoY_Change  DH_YoY_Ch_Per  Total_Inst_Hrs
## Length:81      Length:81      Min.    : 504.6
## Class :character  Class :character  1st Qu.:1937.3
## Mode  :character  Mode  :character  Median :2203.2
##                      Mean    :2165.7
##                      3rd Qu.:2446.8
##                      Max.    :3084.1
##
## Total_Inst_Hrs_Prev_Year  Inst_Hrs_YoY_Change  Total_Inst_Hrs_YoY_Change_Per2
## Length:81                Length:81          Length:81
## Class :character          Class :character    Class :character
## Mode  :character          Mode  :character    Mode  :character
##
##
##
## Cons_Sent      NJURN      RPM      CPIUrban
## Min.    : 70.30  Min.    : 2.900  Min.    : 2908236  Min.    :234.7
## 1st Qu.: 89.00  1st Qu.: 4.100  1st Qu.: 68459347  1st Qu.:241.2
## Median : 93.80  Median : 4.900  Median : 77115921  Median :250.8
## Mean    : 91.49  Mean    : 5.615  Mean    : 70822495  Mean    :250.2
```



```
## 3rd Qu.: 97.90    3rd Qu.: 6.200    3rd Qu.: 85326186    3rd Qu.:257.4
## Max.    :101.40    Max.    :16.600    Max.    :101794185    Max.    :274.1
##                                     NA's    :1
## CPIMedian
## Min.    :0.9755
## 1st Qu.:2.1551
## Median :2.5922
## Mean    :2.5862
## 3rd Qu.:2.9557
## Max.    :5.5690
##
```

```
str(Tng_Ctr_Hour2)
```

```
## Classes 'data.table' and 'data.frame': 81 obs. of 16 variables:
## $ Year : Factor w/ 81 levels "2015-01","2015-02",...: 1 2 3 4 5 6 7 8 9 10
## $ Quarter : Ord.factor w/ 4 levels "Q1"<"Q2"<"Q3"<...: 1 1 1 2 2 2 3 3 3 4 ...
## $ Month : Ord.factor w/ 12 levels "Apr"<"Aug"<"Dec"<...: 5 4 8 1 9 7 6 2 12
## $ Device_Hrs : num 960 944 1429 1097 916 ...
## $ DH_Prev_Year : chr "NA" "NA" "NA" "NA" ...
## $ DH_YoY_Change : chr "NA" "NA" "NA" "NA" ...
## $ DH_YoY_Ch_Per : chr "NA" "NA" "NA" "NA" ...
## $ Total_Inst_Hrs : num 1701 1614 2533 2152 1695 ...
## $ Total_Inst_Hrs_Prev_Year : chr "NA" "NA" "NA" "NA" ...
## $ Inst_Hrs_YoY_Change : chr "NA" "NA" "NA" "NA" ...
## $ Total_Inst_Hrs_YoY_Change_Per2: chr "NA" "NA" "NA" "NA" ...
## $ Cons_Sent : num 98.1 95.4 93 95.9 90.7 96.1 93.1 91.9 87.2 90 ...
## $ NJURN : num 6.8 6.7 6.3 5.8 6 5.9 6.2 5.5 5.1 4.9 ...
## $ RPM : num 65975447 59784666 75751609 73090871 78002935 ...
## $ CPIUrban : num 235 235 236 236 237 ...
## $ CPIMedian : num 1.95 1.95 2.43 2.96 2.5 ...
## - attr(*, ".internal.selfref")=<externalptr>
```

```
Mach1 = subset(Tng_Ctr_Hour2, select = c(Month, Device_Hrs, Cons_Sent, NJURN, RPM, CPIUrban, CPIMedian))
```

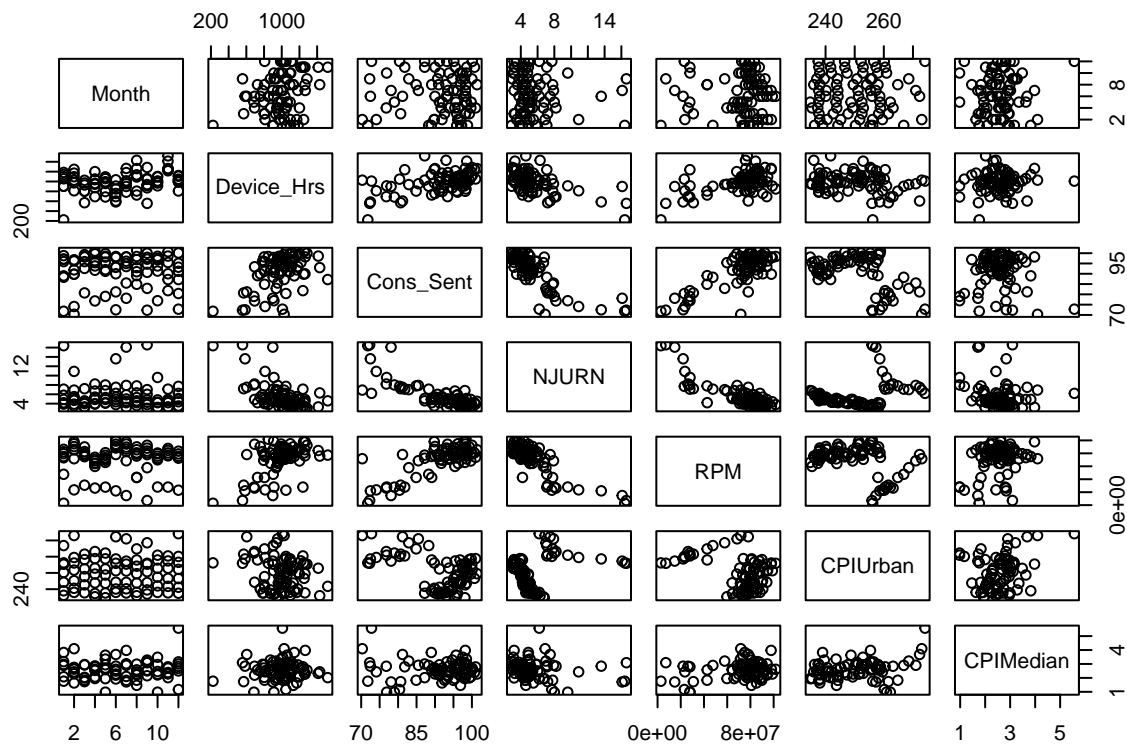
```
LM_Mach = lm(Device_Hrs ~ ., Mach1)
summary(LM_Mach)
```

```
##
## Call:
## lm(formula = Device_Hrs ~ ., data = Mach1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -377.13  -74.78    2.32   71.19  347.28
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  3.403e+02  7.492e+02   0.454  0.651282
## Month.L      2.047e+02  5.664e+01   3.613  0.000602 ***
## Month.Q      1.537e+02  5.915e+01   2.599  0.011634 *
## Month.C     -8.334e+01  5.726e+01  -1.455  0.150511
```

```
## Month^4      -1.169e+02  5.960e+01  -1.962  0.054223 .
## Month^5      -7.904e+01  6.044e+01  -1.308  0.195718
## Month^6      -9.531e+01  7.399e+01  -1.288  0.202428
## Month^7      -1.183e+02  5.881e+01  -2.011  0.048622 *
## Month^8      -2.306e+02  5.606e+01  -4.113  0.000115 ***
## Month^9       1.546e+02  5.682e+01   2.722  0.008389 **
## Month^10     4.619e+01  6.022e+01   0.767  0.445946
## Month^11     -8.832e+01  5.409e+01  -1.633  0.107486
## Cons_Sent     1.811e+00  4.672e+00   0.388  0.699692
## NJURN        -4.201e+00  1.609e+01  -0.261  0.794845
## RPM           4.709e-06  2.417e-06   1.949  0.055773 .
## CPIUrban      9.757e-01  1.984e+00   0.492  0.624510
## CPIMedian     -2.655e+01  2.998e+01  -0.886  0.379123
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 138.6 on 63 degrees of freedom
## (1 observation deleted due to missingness)
## Multiple R-squared:  0.6611, Adjusted R-squared:  0.575
## F-statistic:  7.68 on 16 and 63 DF, p-value: 1.361e-09
```

We show a relatively strong correlation between the model and Device Hours


```
pairs(Mach1)
```





We see some correlation between Consumer Sentiment and NJ Unemployment, and Consumer Sentiment and Revenue Passenger Miles. Its possible there is some redundancy. The CPI Urban and CPI Median don't


show a very strong correlation because CPI Urban is the CPI level, while CPI Median is the percentage change. That is why we're including both.

```
plot(LM_Mach)
```










Our model looks like it is generally normal

Let's look at the residuals


```
res= residuals(LM_Mach)
plot(res)
```



There doesn't appear to be any pattern, which leads us to believe the model is not missing any factors.

Lets take other looks at the residuals

```
hist(res)
```



This looks skewed left, so not completely normal. Maybe we can get a stronger model by removing some independent variables and simplifying the model.

Stepwise

Because we have multiple independent variables, it is important to run a stepwise to see if we can get a stronger model.

```
Mach.Step = step(LM_Mach)
```

```
## Start:  AIC=803.91
## Device_Hrs ~ Month + Cons_Sent + NJURN + RPM + CPIUrban + CPIMedian
##
##           Df Sum of Sq    RSS    AIC
## - NJURN      1      1309 1210999 801.99
## - Cons_Sent   1      2883 1212572 802.10
## - CPIUrban    1      4646 1214335 802.21
## - CPIMedian   1     15064 1224754 802.90
## <none>                 1209689 803.91
## - RPM         1      72925 1282614 806.59
## - Month       11     1127124 2336813 834.58
##
## Step:  AIC=801.99
## Device_Hrs ~ Month + Cons_Sent + RPM + CPIUrban + CPIMedian
```



```
##
##           Df Sum of Sq      RSS      AIC
## - Cons_Sent  1       3536 1214534 800.23
## - CPIUrban   1       7381 1218380 800.48
## - CPIMedian  1      15320 1226319 801.00
## <none>                        1210999 801.99
## - RPM        1      181252 1392251 811.15
## - Month      11     1280870 2491868 837.72
##
## Step: AIC=800.23
## Device_Hrs ~ Month + RPM + CPIUrban + CPIMedian
##
##           Df Sum of Sq      RSS      AIC
## - CPIUrban   1       6674 1221208 798.67
## - CPIMedian  1      20309 1234843 799.55
## <none>                        1214534 800.23
## - Month      11     1353176 2567710 838.12
## - RPM        1      897519 2112053 842.49
##
## Step: AIC=798.67
## Device_Hrs ~ Month + RPM + CPIMedian
##
##           Df Sum of Sq      RSS      AIC
## - CPIMedian  1      14905 1236113 797.64
## <none>                        1221208 798.67
## - Month      11     1346804 2568013 836.13
## - RPM        1     1084290 2305499 847.50
##
## Step: AIC=797.64
## Device_Hrs ~ Month + RPM
##
##           Df Sum of Sq      RSS      AIC
## <none>                        1236113 797.64
## - Month 11     1355720 2591834 834.87
## - RPM   1      1112291 2348404 846.98
```

```
Mach.Step
```

```
##
## Call:
## lm(formula = Device_Hrs ~ Month + RPM, data = Mach1)
##
## Coefficients:
## (Intercept)      Month.L      Month.Q      Month.C      Month^4      Month^5
##  6.146e+02    2.129e+02    1.539e+02   -7.292e+01   -1.308e+02   -8.675e+01
##      Month^6      Month^7      Month^8      Month^9      Month^10      Month^11
## -8.133e+01  -1.351e+02  -2.265e+02    1.634e+02    6.577e+01   -8.436e+01
##           RPM
##  5.342e-06
```

```
summary(Mach.Step)
```

```
##
```


```
## Call:
## lm(formula = Device_Hrs ~ Month + RPM, data = Mach1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -381.44  -76.77    8.75   68.32  321.01
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  6.146e+02  5.103e+01  12.045  < 2e-16 ***
## Month.L      2.129e+02  5.354e+01   3.977  0.000174 ***
## Month.Q      1.539e+02  5.275e+01   2.918  0.004789 **
## Month.C     -7.292e+01  5.296e+01  -1.377  0.173118
## Month^4     -1.308e+02  5.343e+01  -2.448  0.017001 *
## Month^5     -8.675e+01  5.296e+01  -1.638  0.106103
## Month^6     -8.133e+01  5.364e+01  -1.516  0.134196
## Month^7     -1.351e+02  5.339e+01  -2.530  0.013759 *
## Month^8     -2.265e+02  5.353e+01  -4.232  7.2e-05 ***
## Month^9      1.634e+02  5.264e+01   3.104  0.002800 **
## Month^10     6.577e+01  5.168e+01   1.272  0.207602
## Month^11    -8.436e+01  5.141e+01  -1.641  0.105541
## RPM          5.342e-06  6.880e-07   7.765  6.4e-11 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 135.8 on 67 degrees of freedom
## (1 observation deleted due to missingness)
## Multiple R-squared:  0.6537, Adjusted R-squared:  0.5917
## F-statistic: 10.54 on 12 and 67 DF, p-value: 2.594e-11
```

Stepwise reveals that the best way to predict training hours is to use the month and the Revenue Passenger Miles. The other variables were not strong predictors.

This model has a decently strong Adjusted R2, a high F-statistic, and a very low p-value.


Lets look at the residuals of the new model.

```
res.step= residuals(Mach.Step)
plot(res.step)
```



This also looks normally distributed.


```
hist(res.step)
```





this looks to be more normally distributed, though it still slightly skewed to the right.


Final residual check.

```
plot(Mach.Step)
```










Normal QQ plot shows that the model is relatively normal.

```
p1= ggplot(Mach.Step) +
  aes(x=.fitted, y=.resid)+
  geom_point() + geom_abline(intercept = 0, slope = 0, color="red") + geom_smooth(method = "loess", se = FALSE)
labs(title= "Residual Plot")
p1

## 'geom_smooth()' using formula 'y ~ x'
```


Residual Plot



We show some skewedness at the right side of the graph, but generally it appears strong..

Lets look at a comparison of the big model vs the small model.

```
anova(LM_Mach,Mach.Step)
```


```
## Analysis of Variance Table
##
## Model 1: Device_Hrs ~ Month + Cons_Sent + NJURN + RPM + CPIUrban + CPIMedian
## Model 2: Device_Hrs ~ Month + RPM
##   Res.Df    RSS Df Sum of Sq   F Pr(>F)
## 1      63 1209689
## 2      67 1236113 -4    -26424 0.344 0.8472
```

```
anova(Mach.Step,LM_Mach)
```

```
## Analysis of Variance Table
##
## Model 1: Device_Hrs ~ Month + RPM
## Model 2: Device_Hrs ~ Month + Cons_Sent + NJURN + RPM + CPIUrban + CPIMedian
##   Res.Df    RSS Df Sum of Sq   F Pr(>F)
## 1      67 1236113
## 2      63 1209689  4     26424 0.344 0.8472
```

ARIMA Testing

```
plot(FD_train)
```



```
adf.test(FD_train)
```

```
## Warning in adf.test(FD_train): p-value smaller than printed p-value
```

```
##  
## Augmented Dickey-Fuller Test  
##  
## data: FD_train  
## Dickey-Fuller = -5.3294, Lag order = 3, p-value = 0.01  
## alternative hypothesis: stationary
```

P-Value is less than 0.01

Kipps test says differences is required if p-value is < 0.05

```
kpss.test(FD_train)
```

```
## Warning in kpss.test(FD_train): p-value greater than printed p-value
```

```
##
```

```
## KPSS Test for Level Stationarity
```

```
##
```

```
## data: FD_train
```

```
## KPSS Level = 0.16197, Truncation lag parameter = 3, p-value = 0.1
```

P-Value less than 0.05 so no difference needed

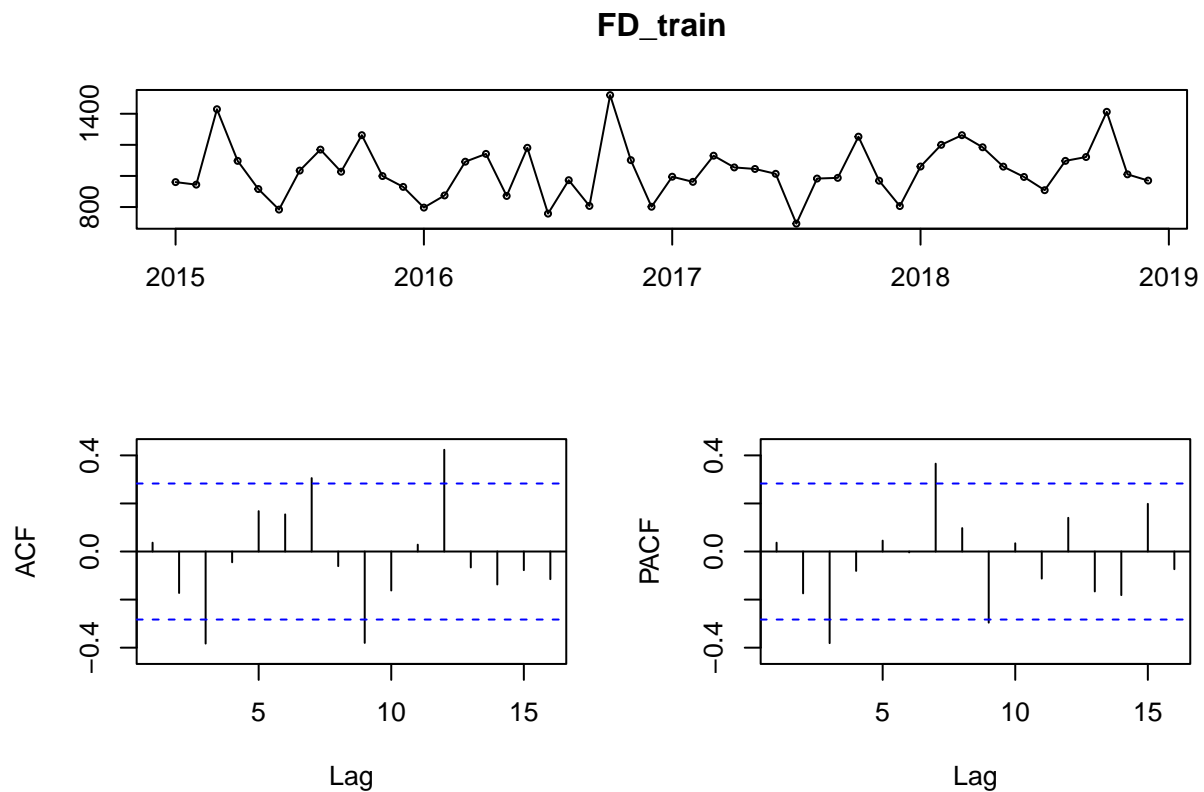
```
nsdiffs(FD_train)
```

```
## [1] 1
```

```
MachDiff = ndiffs(FD_train)
```

```
#tsdisplay plots ACF,PACF and timeseries plot together. How cool!
```


```
tsdisplay(FD_train)
```



We see some significance in the ACF for months 3, 7, 9, and 12. This leads us to believe that there is some dependence. There is significance in the PACF for months 3, 7, and 9.

```
MachDiff2 <- diff(FD_train, differences=1)
```

```
tsdisplay(MachDiff2)
```



There is even ACF significance for months 1, 3, 7, 9, and 12. PACF negative in 1, 3, 4, 6, 8, and 14

```
ndiffs(MachDiff2)
```

```
## [1] 0
```

There is no longer significance

```
ndiffs(MachDiff)
```

```
## [1] 0
```

No significance

```
adf.test(MachDiff2)
```

```
## Warning in adf.test(MachDiff2): p-value smaller than printed p-value
```

```
##
## Augmented Dickey-Fuller Test
##
## data: MachDiff2
## Dickey-Fuller = -7.117, Lag order = 3, p-value = 0.01
## alternative hypothesis: stationary
```

get the partial autocorrelation values

```
Pacf(FD_train, lag.max=20,plot=FALSE)
```

```
##
## Partial autocorrelations of series 'FD_train', by lag
##
##      1      2      3      4      5      6      7      8      9     10     11
## 0.037 -0.174 -0.381 -0.081  0.045 -0.003  0.365  0.097 -0.296  0.034 -0.113
##      12     13     14     15     16     17     18     19     20
## 0.140 -0.166 -0.181  0.198 -0.074  0.001  0.143  0.027 -0.046
```

```
auto.arima(FD_train)
```

```
## Series: FD_train
## ARIMA(0,1,1)(1,1,0)[12]
##
## Coefficients:
##          ma1      sar1
##      -0.8729  -0.3746
## s.e.   0.0903   0.1942
##
## sigma^2 estimated as 27192: log likelihood=-229.02
## AIC=464.03   AICc=464.8   BIC=468.7
```

```
auto.arima(MachDiff)
```

```
## Series: MachDiff
## ARIMA(0,0,0) with non-zero mean
##
## Coefficients:
## intercept
##          0
##
## sigma^2 estimated as 0: log likelihood=Inf
## AIC=-Inf   AICc=-Inf   BIC=-Inf
```

```
auto.arima(MachDiff2)
```

```
## Series: MachDiff2
## ARIMA(0,0,1)(1,1,0)[12]
##
```

```
## Coefficients:
##          ma1      sar1
##      -0.8729 -0.3746
## s.e.   0.0903  0.1942
##
## sigma^2 estimated as 27192:  log likelihood=-229.02
## AIC=464.03   AICc=464.8   BIC=468.7

# or save the model. BIC and AIC is also given as values.
auto_fit <- auto.arima(FD_train, trace=TRUE, stepwise=FALSE)
```

```
##
## ARIMA(0,1,0)(0,1,0)[12] : 484.0533
## ARIMA(0,1,0)(0,1,1)[12] : Inf
## ARIMA(0,1,0)(1,1,0)[12] : 481.0162
## ARIMA(0,1,0)(1,1,1)[12] : Inf
## ARIMA(0,1,1)(0,1,0)[12] : 465.4676
## ARIMA(0,1,1)(0,1,1)[12] : Inf
## ARIMA(0,1,1)(1,1,0)[12] : 464.8046
## ARIMA(0,1,1)(1,1,1)[12] : Inf
## ARIMA(0,1,2)(0,1,0)[12] : 467.7442
## ARIMA(0,1,2)(0,1,1)[12] : Inf
## ARIMA(0,1,2)(1,1,0)[12] : 467.3637
## ARIMA(0,1,2)(1,1,1)[12] : Inf
## ARIMA(0,1,3)(0,1,0)[12] : 470.1542
## ARIMA(0,1,3)(0,1,1)[12] : Inf
## ARIMA(0,1,3)(1,1,0)[12] : 469.8397
## ARIMA(0,1,3)(1,1,1)[12] : Inf
## ARIMA(0,1,4)(0,1,0)[12] : 466.5286
## ARIMA(0,1,4)(0,1,1)[12] : Inf
## ARIMA(0,1,4)(1,1,0)[12] : 466.677
## ARIMA(0,1,5)(0,1,0)[12] : Inf
## ARIMA(1,1,0)(0,1,0)[12] : 475.7162
## ARIMA(1,1,0)(0,1,1)[12] : Inf
## ARIMA(1,1,0)(1,1,0)[12] : 475.5616
## ARIMA(1,1,0)(1,1,1)[12] : Inf
## ARIMA(1,1,1)(0,1,0)[12] : 467.7553
## ARIMA(1,1,1)(0,1,1)[12] : Inf
## ARIMA(1,1,1)(1,1,0)[12] : Inf
## ARIMA(1,1,1)(1,1,1)[12] : Inf
## ARIMA(1,1,2)(0,1,0)[12] : 469.5689
## ARIMA(1,1,2)(0,1,1)[12] : Inf
## ARIMA(1,1,2)(1,1,0)[12] : 469.8107
## ARIMA(1,1,2)(1,1,1)[12] : Inf
## ARIMA(1,1,3)(0,1,0)[12] : 470.9884
## ARIMA(1,1,3)(0,1,1)[12] : Inf
## ARIMA(1,1,3)(1,1,0)[12] : 472.2383
## ARIMA(1,1,4)(0,1,0)[12] : Inf
## ARIMA(2,1,0)(0,1,0)[12] : 476.926
## ARIMA(2,1,0)(0,1,1)[12] : Inf
## ARIMA(2,1,0)(1,1,0)[12] : 476.5811
## ARIMA(2,1,0)(1,1,1)[12] : Inf
## ARIMA(2,1,1)(0,1,0)[12] : 470.298
## ARIMA(2,1,1)(0,1,1)[12] : Inf
```

```
## ARIMA(2,1,1)(1,1,0)[12] : 469.8879
## ARIMA(2,1,1)(1,1,1)[12] : Inf
## ARIMA(2,1,2)(0,1,0)[12] : 472.0187
## ARIMA(2,1,2)(0,1,1)[12] : Inf
## ARIMA(2,1,2)(1,1,0)[12] : 471.8468
## ARIMA(2,1,3)(0,1,0)[12] : Inf
## ARIMA(3,1,0)(0,1,0)[12] : 468.4805
## ARIMA(3,1,0)(0,1,1)[12] : Inf
## ARIMA(3,1,0)(1,1,0)[12] : 467.4776
## ARIMA(3,1,0)(1,1,1)[12] : Inf
## ARIMA(3,1,1)(0,1,0)[12] : 468.2049
## ARIMA(3,1,1)(0,1,1)[12] : Inf
## ARIMA(3,1,1)(1,1,0)[12] : 468.4294
## ARIMA(3,1,2)(0,1,0)[12] : 470.7111
## ARIMA(4,1,0)(0,1,0)[12] : 470.3022
## ARIMA(4,1,0)(0,1,1)[12] : Inf
## ARIMA(4,1,0)(1,1,0)[12] : 469.8215
## ARIMA(4,1,1)(0,1,0)[12] : 470.7823
## ARIMA(5,1,0)(0,1,0)[12] : 472.0753
##
##
## Best model: ARIMA(0,1,1)(1,1,0)[12]
```

```
auto_fit
```


```
## Series: FD_train
## ARIMA(0,1,1)(1,1,0)[12]
##
## Coefficients:
##          ma1      sar1
##      -0.8729 -0.3746
## s.e.   0.0903   0.1942
##
## sigma^2 estimated as 27192: log likelihood=-229.02
## AIC=464.03 AICc=464.8 BIC=468.7
```

```
forecast(auto_fit,h=5,level=c(99.5))
```

```
##          Point Forecast  Lo 99.5  Hi 99.5
## Jan 2019      1113.566  650.6612 1576.470
## Feb 2019      1188.517  721.8870 1655.148
## Mar 2019      1290.320  819.9933 1760.647
## Apr 2019      1213.371  739.3761 1687.365
## May 2019      1131.839  654.2051 1609.473
```


```
plot(forecast(auto_fit,h=5,level=c(99.5)))
```

Forecasts from ARIMA(0,1,1)(1,1,0)[12]




```
Acf(auto_fit$residuals)
```


Series auto_fit\$residuals




```
hist(auto_fit$residuals)
```

Histogram of auto_fit\$residuals




Appears to be skewed left

```
tsdiag(auto_fit)
```




Seasonal ARIMA

```
tsdisplay(diff(FD_train,4))
```



```
# It did not make it stationary. So lets take lag 1 diffence  
tsdisplay(diff(diff(FD_train,4)))
```



```
ndiffs(diff(FD_train,12))
```

```
## [1] 1
```

Checking the model

```
SeasMach = diff(diff(FD_train,12))
```

```
Pacf(SeasMach)
```

Series SeasMach

